

Support Vector Regression (SVR) for regression


Freddy Hernández

6 de noviembre de 2019



SVR were developed by Drucker et al. (1996)

Url: <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.21.5909>



Support Vector Regression Machines

Harris Drucker* Chris J.C. Burges** Linda Kaufman**
Alex Smola** Vladimir Vapnik *

*Bell Labs and Monmouth University
Department of Electronic Engineering
West Long Branch, NJ 07764

**Bell Labs *AT&T Labs

Abstract


A new regression technique based on Vapnik's concept of support vectors is introduced. We compare support vector regression (SVR) with a committee regression technique (bagging) based on regression trees and ridge regression done in feature space. On the basis of these experiments, it is expected that SVR will have advantages in high dimensionality space because SVR optimization does not depend on the dimensionality of the input space.

An overview of Support Vector Machines

<https://www.svm-tutorial.com/2017/02/svms-overview-support-vector-machines/>



The objective of the Support Vector Machine algorithm is to find a margin of tolerance (ϵ) to minimize error, individualizing the hyperplane which maximizes the margin, keeping in mind that part of the error is tolerated.



- **Minimize:**

$$\frac{1}{2} \|w\|^2 + C \sum_{i=1}^N (\xi_i + \xi_i^*)$$

- **Constraints:**

$$y_i - wx_i - b \leq \epsilon + \xi_i$$

$$wx_i + b - y_i \leq \epsilon + \xi_i^*$$

$$\xi_i, \xi_i^* \geq 0$$



The package **e1071** contains the `svm` function used for SVMs. To install the package:

```
install.packages("e1071")
```

To load the package:

```
library(e1071)
```

The main function is:

```
svm(formula, data, scale=TRUE,  
     kernel=linear or polynomial or radial or sigmoid,  
     degree=3, gamma=1/n, coef0=0, cost=1, ...)
```


Consult the vignette:

<https://cran.r-project.org/web/packages/e1071/vignettes/svmdoc.pdf>



Example with simulated data

```
set.seed(1234)
x <- sort(runif(n=40, min=0, max=5)) # sort for convenience
set.seed(1234)
y <- sin(x) + rnorm(40, sd=0.3)
plot(x, y, pch=20, las=1)
```



This function can be used to obtain the MSE.

```
mse <- function(y, y_hat) mean((y - y_hat)^2)
```



```
mod_lin <- svm(y ~ x, kernel="linear")
```

To obtain \hat{y} .

```
y_hat_lin <- predict(mod_lin)
```

To obtain the correlation coefficient and MSE.

```
cor(y, y_hat_lin)
```

```
## [1] 0.7852445
```


```
mse(y, y_hat_lin)
```

```
## [1] 0.2687447
```



Linear kernel

```
plot(x, y, pch=20)  
points(x=x, y=y_hat_lin, type="l", lwd=2, col="red")
```



Polynomial kernel

```
mod_pol <- svm(y ~ x, kernel="polynomial")
```

To obtain \hat{y} .

```
y_hat_pol <- predict(mod_pol)
```

To obtain the correlation coefficient and MSE.

```
cor(y, y_hat_pol)
```

```
## [1] 0.6468922
```


```
mse(y, y_hat_pol)
```

```
## [1] 0.4344635
```



Polynomial kernel

```
plot(x, y, pch=20)  
points(x=x, y=y_hat_pol, type="l", lwd=2, col="blue")
```



```
mod_rad <- svm(y ~ x, kernel="radial")
```

To obtain \hat{y} .

```
y_hat_rad <- predict(mod_rad)
```

To obtain the correlation coefficient and MSE.

```
cor(y, y_hat_rad)
```


```
## [1] 0.9506258
```

```
mse(y, y_hat_rad)
```


```
## [1] 0.06821404
```



```
plot(x, y, pch=20)  
points(x=x, y=y_hat_rad, type="l", lwd=2, col="forestgreen")
```



Comparisons



Below we have some .gif figures to explore the tune parameter effects.

- https://github.com/rdaymedellin/tutoriales_Rday_2019/blob/master/Machine%20Learning%20SVM/animation_svm_pol_reg.gif
- https://github.com/rdaymedellin/tutoriales_Rday_2019/blob/master/Machine%20Learning%20SVM/animation_svm_rad_reg.gif



We can use the `tune.svm` function.

```
tune.svm(formula, data, kernel,  
          degree, gamma, coef0, cost)
```

- the classification error is used for classification.
- the mean squared error for regression.



Tuning parameters for polynomial

```
pol_tune = tune.svm(y~x, kernel="polynomial",
                    degree=c(2, 3, 4),
                    gamma=c(0.1, 1, 2),
                    coef0=c(0.1, 0.5, 1, 2, 3))
summary(pol_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   degree gamma coef0
##       4      1      2
##
## - best performance: 0.08341557
##
## - Detailed performance results:
```

	degree	gamma	coef0	error	dispersion
## 1	2	0.1	0.1	0.32067808	0.21307867
## 2	3	0.1	0.1	0.65578295	0.45230278
## 3	4	0.1	0.1	0.78137035	0.54532617
## 4	2	1.0	0.1	0.16156050	0.15002103
## 5	3	1.0	0.1	0.32109973	0.35392559
## 6	4	1.0	0.1	0.71894997	1.10410176
## 7	2	2.0	0.1	0.15656789	0.15172337
## 8	3	2.0	0.1	0.24171162	0.25211007
## 9	4	2.0	0.1	0.54251151	0.27222527

Tuning parameters for radial

```
rad_tune = tune.svm(y~x, kernel="radial",  
                    gamma=c(0.1, 0.5, 1, 1.5, 2))  
summary(rad_tune)
```

```
##  
## Parameter tuning of 'svm':  
##  
## - sampling method: 10-fold cross validation  
##  
## - best parameters:  
##   gamma  
##     0.5  
##  
## - best performance: 0.09273725  
##  
## - Detailed performance results:  
##   gamma      error dispersion  
## 1    0.1 0.16878872 0.12622542  
## 2    0.5 0.09273725 0.08338983  
## 3    1.0 0.09790368 0.08082934  
## 4    1.5 0.10448328 0.07865747  
## 5    2.0 0.10722114 0.07893561
```



```
best_pol <- pol_tune$best.model  
y1 <- predict(best_pol)  
cor(y, y1)
```

```
## [1] 0.9512053
```

```
mse(y, y1)
```

```
## [1] 0.06545502
```

```
best_rad <- rad_tune$best.model  
y2 <- predict(best_rad)  
cor(y, y2)
```

```
## [1] 0.9468661
```


```
mse(y, y2)
```

```
## [1] 0.07281149
```



Comparisons

Default



Best tune parameters

